



AI Enterprise Workflow Study Group

Course 6, Week 2

5/23/2020

Course & Study Group Schedule

AI Enterprise Workflow Study Group		
Session	Topic	Date
Overview Webinar	Webinar with instructor, Ray Lopez	15-Feb
Course 1 Week 1	Course intro	22-Feb
Course 1 Week 2	Data ingestion, cleaning, parsing, assembly	29-Feb
Course 2 Week 1	Exploratory data analysis & visualization	7-Mar
Course 2 Week 2	Estimation and NHT	14-Mar
Course 3 Week 1	Data transformation and feature engineering	21-Mar
Course 3 Week 2	Pattern recognition and data mining best practices	28-Mar
Course 4 Week 1	Model evaluation and performance metrics	18-Apr
Course 4 Week 2	Building machine learning and deep learning models	25-Apr
Course 5 Week 1	Deploying models	2-May
Course 5 Week 2	Deploying models using Spark	9-May
Course 6 Week 1	Feedback loops and monitoring	16-May
Course 6 Week 2	Hands on with OpenScale and Kubernetes	23-May
Course 6 Week 3	Captstone project week 1	30-May
Course 6 Week 4	Captstone project week 2	6-Jun

Agenda

- Check in
- Discussion
- Next steps

Course 6 Week 2 learning objectives

- Explain the management and diagnostic services available through Watson OpenScale
- Implement the basic tooling required to monitor model performance
- Deploy a model and manage containers with Kubernetes

Watson Machine Learning

Watson Machine Learning is a service on IBM Cloud with features for training and deploying machine learning models and neural networks. It is based on Kubernetes and Docker components.

Programming Interfaces:

- Python client library
- Command line interface
- REST API external link

Training infrastructure:

- Distributed training to run training jobs across multiple servers
- Popular frameworks are supported so you can work with your preferred libraries
- GPUs for faster training
- Hyperparameter optimization for training complex neural networks

Deployment infrastructure:

- Deploy trained models - as a web service, for batch processing, or with streaming data
- Deploy Python functions to simplify AI solutions

IBM Watson OpenScale

IBM Watson OpenScale is the open solution for businesses to operationalize trusted AI and extend their deployments enterprise-wide.


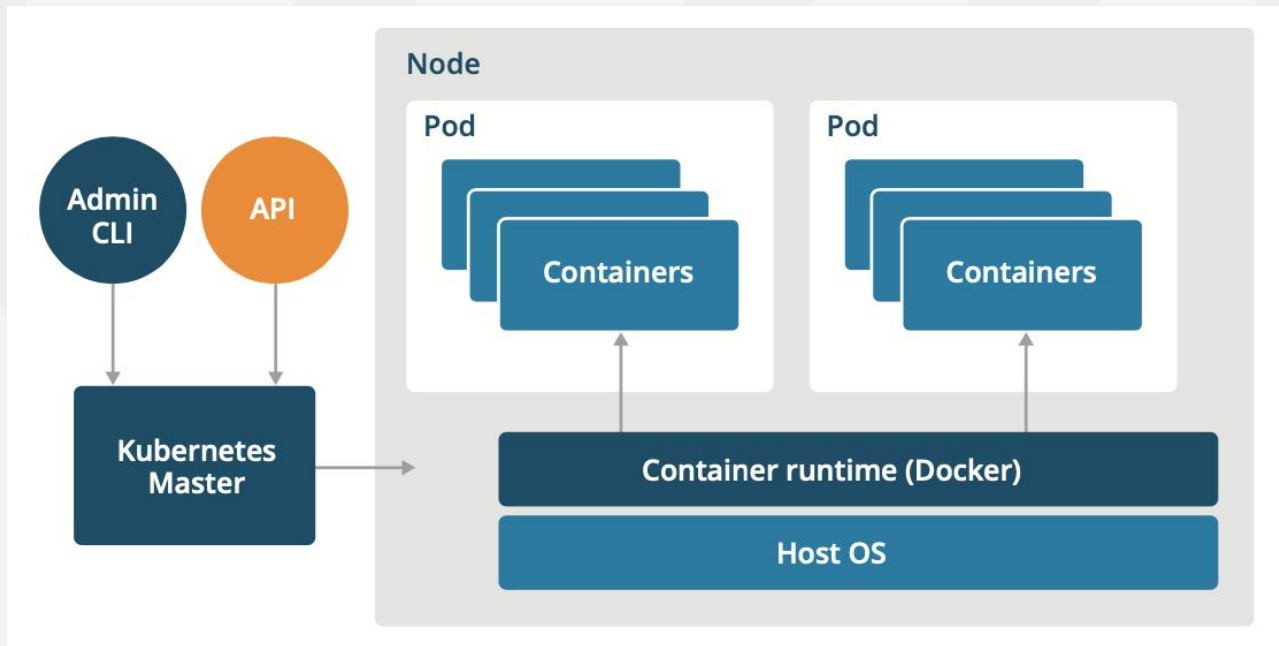
Use Cases:

- Performance Monitoring
- Quality Monitoring
- Fairness Monitoring
- Explainability
- Drift monitoring

Four Main Areas in Tool:

- Insights (Dashboard)
- Explanations
- Configuration
- Support

Kubernetes



The logo for Twiml, featuring the word "twiml" in a white, lowercase, sans-serif font. A small blue horizontal bar is positioned above the "i".

twiml